

CHALMERS
UNIVERSITY OF TECHNOLOGY

Statistical methods in Data Science and AI

Marina Axelson-Fisk
30 september, 2019



Inference in graphical models


Given a graphical model, we want to answer questions of interest.

- **Marginal inference:** what is the marginal probability of a given variable Y in our graph, summing out the rest?

$$P(Y = y) = \sum_{x_1} \sum_{x_2} \dots \sum_{x_n} P(Y = y, X_1 = x_1, X_2 = x_2, \dots, X_n = x_n)$$

- **Maximum a posteriori (MAP) inference:** what is the most likely assignment to the variables in the graph (possibly conditioned on data)?

$$\max_{x_1, \dots, x_n} P(Y = y, X_1 = x_1, \dots, X_n = x_n)$$




Inference algorithms in graphical models

Exact inference

- Variable elimination
- Message passing/belief propagation
- Junction trees


Approximative inference

- Stochastic simulation
- Markov chain Monte Carlo (MCMC)
- Variational algorithms



Example: variable elimination in a chain graph

Random variables: A, B, C, D, E



each taking n possible values $\{1, 2, \dots, n\}$, then the **marginal probability** of E

$$P(E = e) = \sum_{a=1}^n \sum_{b=1}^n \sum_{c=1}^n \sum_{d=1}^n P(A = a, B = b, C = c, D = d, E = e) = \sum_{a,b,c,d} P(a, b, c, d)$$

we can utilize the chain structure to reduce the number of operations by **variable elimination**.

Example: variable elimination in a chain graph

Exploit the structure "inside-out" or from "leaf-to-top"


$$\begin{aligned}
 P(E = e) &= \sum_a \sum_b \sum_c \sum_d P(A = a, B = b, C = c, D = d, E = e) \\
 &= \sum_a \sum_b \sum_c \sum_d P(a)P(b|a)P(c|b)P(d|c)P(e|d) \\
 &= \sum_b \sum_c \sum_d P(c|b)P(d|c)P(e|d) \sum_a P(b|a)P(a) \\
 &= \sum_b \sum_c \sum_d P(c|b)P(d|c)P(e|d) P(b) = \dots = \\
 &= \sum_d P(e|d)P(d) = P(e)
 \end{aligned}$$

Message passing/belief propagation


- Variable elimination can be seen as "passing a message" (information), or "**propagating a belief**" from one node to the next.
- This is the basic framework for computing various entities in **Hidden Markov Models** (HMMs) and **Linear dynamical systems** (LDS)
- In continuous distributions $p(x|\theta)$ message passing corresponds to passing on **parameter values** θ between neighboring nodes.

Inference on trees


- Each node sends out the **product** of the messages received **from the parents to the children**
- I.e. message passing is an abstract notion of **conditional (in)dependence**



undirected tree




directed tree




polytree

Inference on trees

- The *junction tree algorithm* (or *clique tree algorithm*) is a generalization of message passing to arbitrary graphs




Module 4.1: Kernel methods




Kernel methods: motivation

- Given a training set $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$
 - y_i response
 - \mathbf{x}_i feature vector
- there are numerous tools for detecting linear relations
 - Ridge regression
 - Support vector machines (SVMs)
 - Principal component analysis (PCA)
- But what if the relationship is nonlinear?



Motivating example: binary classification

- **Training set:** $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N, y_i \in \{-1, +1\}$
- **Objective:** learn a function $f(\mathbf{x})$ such that $y_i = \text{sign}(f(\mathbf{x}_i))$




Linear classification

When classes are linearly separable, the boundary is a *hyperplane*.

$$\mathbf{w}^T \mathbf{x} + b = 0$$


$$\text{If } \mathbf{w}^T \mathbf{x} + b \begin{cases} > 0 & y = \text{red} \\ < 0 & y = \text{green} \end{cases}$$

But which line should we choose?



Linear classification

- How large margins do we have between the classes?
- How do we maximize that margin?




How do we maximize the margin?

- Let $\mathbf{x}_n \in \mathcal{D}$ be the point **closest** to the hyperplane $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = 0$
- Normalize \mathbf{w} :

$$|\mathbf{w}^T \mathbf{x}_n| = 1 \quad \Rightarrow \text{The canonical hyperplane}$$
- The distance between \mathbf{x}_n and the plane

$$\text{distance} = \frac{1}{\|\mathbf{w}\|} |\mathbf{w}^T \mathbf{x}_n - \mathbf{w}^T \mathbf{x}| = \frac{1}{\|\mathbf{w}\|}$$
- Maximize this distance.



Support vector machines (SVMs)

For two linearly separable classes with class labels $y_i \in \{-1, +1\}$

- construct *two* supporting hyperplanes, one for each class


$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 \text{ for } y_i = -1$$

$$\mathbf{w}^T \mathbf{x}_i + b \geq +1 \text{ for } y_i = +1$$

The corresponding supporting hyperplanes are thus

$$H_{-1} = \{\mathbf{x}_i: \mathbf{w}^T \mathbf{x}_i + b = -1, y_i = -1\}$$

$$H_{+1} = \{\mathbf{x}_i: \mathbf{w}^T \mathbf{x}_i + b = +1, y_i = +1\}$$



Support vector machines (SVMs)

In two dimensions the separating hyperplanes are lines on the form

$$f_{\mathbf{w}}(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i + b = w_1 x_1 + w_2 x_2 + b = 0$$

Maximizing the margin = minimizing $\|\mathbf{w}\|$
under the constraints


$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 \text{ for } y_i = -1$$

$$\mathbf{w}^T \mathbf{x}_i + b \geq +1 \text{ for } y_i = +1$$

or

$$\mathbf{w}^* = \arg \min_{\mathbf{w}, b} \left\{ \frac{1}{2} \|\mathbf{w}\|^2 : y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \right\}$$

using *Lagrange multipliers*.



Lagrange multipliers

- We want to optimize a function $f(x)$ subject to a constraint $g(x) = 0$.

- We form the **Lagrangian function**


$$L(x, \lambda) = f(x) - \lambda g(x)$$

where λ is a **Lagrange multiplier**.

- We optimize by computing

$$\frac{dL}{dx} = 0 \quad \text{and} \quad \frac{dL}{d\lambda} = 0$$

and solve the corresponding equation system.



Support vector machines (SVMs)

Lagrangian function


$$L(\mathbf{w}, b, \lambda) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \lambda_i \left(y_i (\mathbf{w}^T \mathbf{x}_i - b) \right)$$

- $\lambda = (\lambda_1, \dots, \lambda_N)$: **Lagrange multipliers**
- N : **number of constraints**

$$\frac{dL}{dw_i} = 0, i = 1, \dots, N \text{ and } \frac{dL}{db} = 0$$

$$\Rightarrow \mathbf{w} = \sum_{i=1}^N \lambda_i y_i \mathbf{x}_i \text{ and } \sum_{i=1}^N \lambda_i y_i = 0$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}, b} \left\{ \frac{1}{2} \|\mathbf{w}\|^2 : y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \right\}$$



Kernelized SVMs

Dual representation

Instead we maximize


$$\tilde{L}(\lambda) = \sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$$

under the constraints

$$\lambda_i \geq 0 \quad \text{and} \quad \sum_{i=1}^N \lambda_i y_i = 0$$

where $k(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$ is a **kernel function**.


$$\mathbf{w} = \sum_{i=1}^N \lambda_i y_i \mathbf{x}_i \quad \text{and} \quad \sum_{i=1}^N \lambda_i y_i = 0$$



Kernelized SVMs

To classify a new data point \mathbf{x}^* we observe the sign of

$$f(\mathbf{x}^*) = \sum_{i=1}^N \lambda_i y_i k(\mathbf{x}^*, \mathbf{x}_i) + b$$



Soft margin SVMs

When the training set cannot be perfectly separated we introduce **slack variables**


$$y_i (\mathbf{w}^T \mathbf{x} + b) \geq 1 - \xi_i$$

such that $\xi_i \leq 1$ for points on the correct side.

We want to minimize the misclassification rate, i.e. minimize

$$\sum_{i=1}^N \mathbb{I}\{\xi_i - 1\} \quad \text{where} \quad \mathbb{I}(x) = \begin{cases} 1 & \text{if } x \leq 0 \\ 0 & \text{if } x > 0 \end{cases}$$

NP-complete!



Soft margin SVMs

Use the upper bound

$$\sum_{i=1}^N \mathbb{I}\{\xi_i - 1\} \leq \sum_{i=1}^N \xi_i$$


and minimize

$$C \sum_{i=1}^N \xi_i + \frac{1}{2} \|\mathbf{w}\|^2, \quad C > 0$$

under the **soft margin** constraints

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$$

C is a trade-off between
misclassification and
complexity




Kernelized soft margin SVMs

Lagrangian function

$$L(\mathbf{w}, b, \boldsymbol{\lambda}) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \xi_i - \sum_{i=1}^N \lambda_i \left(y_i (\mathbf{w}^T \mathbf{x}_i - 1 + \xi_i) \right) - \sum_{i=1}^N \mu_i \xi_i$$

with Lagrange multipliers

$$\boldsymbol{\lambda} = (\lambda_1, \dots, \lambda_N), \boldsymbol{\mu} = (\mu_1, \dots, \mu_n)$$




Kernelized soft margin SVMs

The dual form is the same as before


$$\tilde{L}(\lambda) = \sum_{i=1}^N \lambda_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \lambda_i \lambda_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j)$$

but the maximization constraints become limited by C

$$0 \leq \lambda_i \leq C \quad \text{and} \quad \sum_{i=1}^N \lambda_i y_i = 0$$




Nonlinear classification



- There is no linear classifier that can separate **red** from **green**.

Kernel methods: motivation


- **Solution:**
 - map the data into a (possibly high-dimensional) vector space where the relation becomes linear
 - apply the linear algorithm in this space



Nonlinear classification

- There is no linear classifier that can separate **red** from **green**.
- However, the following function can separate the regions perfectly


$$f(x) = x^2 - r = \underbrace{\langle (1, -r), \rangle}_{\mathbf{w}^T} \underbrace{(x^2, 1)}_{\phi(x)}, \text{ for } a < r < b$$



- By mapping x to **feature space** $\phi(x) = (x^2, 1) \in \mathbb{R}^2$ the nonlinear problem has become **linear**.

Nonlinear classification

- There is no linear classifier that can separate **red** from **green**.
- However, a **conic section** separates them perfectly




$$f(\mathbf{x}) = ax_1^2 + bx_1x_2 + cx_2^2 + dx_1 + ex_2 + g$$

$$= \underbrace{\langle (a, b, c, d, e, g) \rangle}_{\mathbf{w}^T} \underbrace{\langle (x_1^2, x_1x_2, x_2^2, x_1, x_2, 1) \rangle}_{\phi(\mathbf{x})}, \quad \phi \in \mathbb{R}^6$$


Kernel methods: motivation

- **Problem:**
 - computationally difficult to represent data in high dimensions



Kernel methods: motivation

- **Alternative:**
 - compute **similarity measure** between vectors in feature space
 - apply algorithms based on similarity measures



Kernel definition

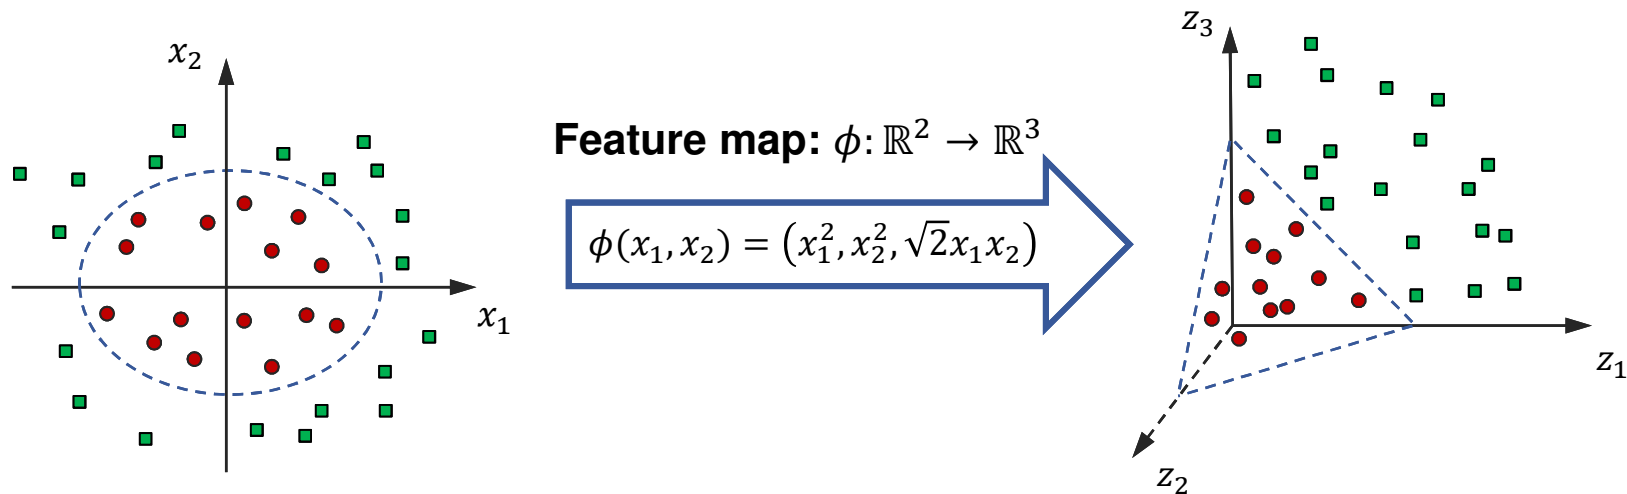
- For input vectors $\mathbf{x}, \mathbf{z} \in \mathcal{X}$ and a mapping $\phi: \mathcal{X} \rightarrow \mathbb{R}^N$

$$k(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})^T \phi(\mathbf{z})$$

is a *kernel function*.

- **The kernel trick:** we **don't** need the coordinates of the data in feature space. Just the **inner product** between vectors.

Example




Kernel: $k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z} = (x_1 z_1 + x_2 z_2)^2 = x_1^2 z_1^2 + 2x_1 z_1 x_2 z_2 + x_2^2 z_2^2 =$
 $= (x_1^2, \sqrt{2}x_1 x_2, x_2^2)(z_1^2, \sqrt{2}z_1 z_2, z_2^2)^T = \phi(\mathbf{x})^T \phi(\mathbf{z})$

where $\phi(\mathbf{x}) = (x_1^2, \sqrt{2}x_1 x_2, x_2^2)^T$ is the nonlinear feature mapping.

Kernel functions

Another view:

- a kernel $k(\mathbf{x}, \mathbf{z})$ is a **measure of similarity** between vectors $\mathbf{x}, \mathbf{z} \in \mathcal{X}$ where \mathcal{X} is some abstract space.
- or simply a **distance measure** between points in feature space




Kernel functions

However, a feature space is **not unique** to a given kernel:

$k(\mathbf{x}, \mathbf{z}) = \langle \mathbf{x}, \mathbf{z} \rangle^2$ is also the kernel that computes the inner product of the map

$$\psi(x_1, x_2) = (x_1^2, x_2^2, x_1x_2, x_2, x_1) \in \mathbb{R}^4$$

Moreover, every prospective kernel does **not** necessarily correspond to a dot product in some space.




Kernel functions

Mercer's condition:

A continuous, symmetric, positive semi-definite kernel function can be written as a dot product of vectors in a higher dimension

$$k(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})^T \phi(\mathbf{z})$$

Positive semi-definite: a symmetric matrix with positive eigenvalues



Let's summarize


- We have a data set $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^N$ where $\mathbf{x} \in \mathbb{R}^D$ and $y_i \in \{-1, +1\}$
- We want a nonlinear projection $\phi(\mathbf{x})$ onto a higher dimension

$$\phi: \mathbb{R}^D \rightarrow \mathbb{R}^{D+d}, \quad d > 0$$

where classes *have a better chance* of being linearly separable.

Cover's theorem (informally):

"A nonlinear projection in a high-dimensional space is more likely to be linearly separable than in a low-dimensional space"



Let's summarize


- The separating hyperplane in \mathbb{R}^{D+d} will be defined by

$$\sum_{j=1}^{D+d} w_j \phi_j(\mathbf{x}) + b = \mathbf{w}^T \phi(\mathbf{x}) + b = 0$$

- The optimal hyperplane is given by


$$\mathbf{w} = \sum_{i=1}^N \lambda_i y_i \phi(\mathbf{x}_i) \Leftrightarrow \sum_{i=1}^N \lambda_i y_i k(\mathbf{x}_i, \mathbf{x}) = 0$$

where $k(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$ is a kernel function and λ_i Lagrange multipliers.



How to choose the mapping $\phi(\mathbf{x})$?


- Choosing an optimal feature space is non-trivial
- The **kernel trick** reduces this to choosing the best kernel, and determine the corresponding implicit mapping $\phi(\mathbf{x})$.
- Performance of the algorithm highly depends on the kernel
- The best kernel depends on the specific problem
- Kernels can be applied to
 - Numeric vectors
 - Strings
 - Trees
 - Graphs



How to choose the best kernel

We want the kernel to be

- **Valid**: an implicit mapping must exist
 - = a kernel that can be expressed as the dot product of two vectors
 - = satisfy the Mercer's condition of positive semi-definiteness
- **Accurate**: embody the "true" similarity between objects
- **Appropriate**: generalize outside training data
- **Efficient**: computations should be feasible



Which kernels meet Mercer's condition?

- **Linear** kernels


$$k(\mathbf{x}, \mathbf{z}) = \mathbf{x}^T \mathbf{z}$$

- **Polynomial** kernels

$$k(\mathbf{x}, \mathbf{z}) = \left(1 + \mathbf{x}^T \mathbf{z}\right)^n$$

- **Radial basis function (RBF) kernels**

$$k(\mathbf{x}, \mathbf{z}) = \exp\left(-\frac{1}{2}\|\mathbf{x} - \mathbf{z}\|^2\right)$$



Radial basis function (RBF) kernels

- The *squared exponential* (SE) or *Gaussian* kernel

$$k(\mathbf{x}, \mathbf{z}) = \exp \left(-\frac{1}{2} (\mathbf{x} - \mathbf{z})^T \Sigma^{-1} (\mathbf{x} - \mathbf{z}) \right)$$


- If the covariance matrix Σ is diagonal, we get

$$k(\mathbf{x}, \mathbf{z}) = \exp \left(-\frac{1}{2} \sum_j \frac{1}{\sigma^2} (x_j - z_j)^2 \right)$$

- If Σ is spherical

$$k(\mathbf{x}, \mathbf{z}) = \exp \left(-\frac{1}{2\sigma^2} \|\mathbf{x} - \mathbf{z}\|^2 \right)$$

An RBF kernel




Kernels for comparing text documents

- If x_{ij} = the number of times word j occurs in document i

$$k(\mathbf{x}, \mathbf{z}) = \frac{\mathbf{x}^T \mathbf{z}}{\|\mathbf{x}\|_2 \|\mathbf{z}\|_2} = \cos(\theta)$$

called the *cosine similarity*.




Matern kernel

- Commonly used in Gaussian processes

$$k(r) = \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}r}{l} \right)^\nu B_\nu \left(\frac{\sqrt{2\nu}r}{l} \right) \rightarrow \text{SE kernel as } \nu \rightarrow \infty$$

where $r = \|\mathbf{x} - \mathbf{z}\|$, $\nu \geq 0$, $l > 0$ and B_ν a modified Bessel function.

$$k(r) = \exp(-r/l) \text{ when } \nu = 1/2$$



String kernels

- Consider two strings \mathbf{x} and \mathbf{z} of lengths $D_{\mathbf{x}}$ and $D_{\mathbf{z}}$, defined on a protein alphabet
 - $\mathcal{A} = \{A, R, N, D, C, E, Q, G, H, I, L, K, M, F, P, S, T, W, Y, V\}$

where

\mathbf{x} ($D_{\mathbf{x}} = 110$):

IPTSALVKETLALLSTHRTLLIANETLRIPVPVHKNHQLCTEEIFQGIGTL
ESQTVQGGTVERLFKNLSLIKKYIDGQKKKCGEERRRVNQFLDY**LQE**FLGV
MNTEWI

\mathbf{z} ($D_{\mathbf{z}} = 153$):

PHRRDLCSRSIWLARKIRSDLTALTESYVKHQGLWSELTEAER**LQEN**LQAY
RTFHVLLARLLEDQQVHFTPTGDFHQAIHTLLLQVAAFAYQIEELMILLE
YKIPRNEADGMLFEKKLWGLKV**LQELS**QWTVRSIHDLRFISSHQTGIP

Similarity measure:
number of common substrings

$$k(\mathbf{x}, \mathbf{z}) = \sum_{s \in \mathcal{A}^*} w_s \phi_s(\mathbf{x}) \phi_s(\mathbf{z})$$

where s is a substring, $w_s \geq 0$
and \mathcal{A}^* the set of all substrings
from \mathcal{A} .

String kernels


- If $w_s = 1$ for a nonempty substring $|s| > 0$:
 $\phi(\mathbf{x})$ = number of times each char in \mathcal{A} occurs in \mathbf{x}
 - **bag-of-characters model**
- If we require each substring s to be surrounded by white space
 $\phi(\mathbf{x})$ = number of times each word s occurs in \mathbf{x}
 - **bag-of-words model**
- If we only consider strings with $|s| = k$ we get the **k -spectrum kernel**

Similarity measure:
number of common substrings

$$k(\mathbf{x}, \mathbf{z}) = \sum_{s \in \mathcal{A}^*} w_s \phi_s(\mathbf{x}) \phi_s(\mathbf{z})$$

where s is a substring, $w_s \geq 0$
and \mathcal{A}^* the set of all substrings from \mathcal{A} .

An SVM as a neural network




Gaussian processes

- **Linear regression:** determine relation f between response y and independent variable x

$$y = f(x) + \epsilon$$


where f is assumed to be linear: $f(x) = \beta_0 + \beta_1 x$

- **Bayesian linear regression:** determine a posterior distribution over the parameters β_0 and β_1 that gets updated whenever new data is available.
- **Gaussian processes:** finds a posterior distribution over the possible **functions** $f(x)$ consistent with the observed data and a suitable prior.




Gaussian processes

- The current example isn't really linear.
- Quadratic function
$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \epsilon$$
- Three parameters to estimate: $\beta_0, \beta_1, \beta_2$
- But what if we don't know how many parameters we should use?
- Instead of searching for suitable parameter values for a fixed number of parameters (and a fixed function), we want to **search among all functions** that fit our data.



Gaussian processes

- We need to define a prior over the function space.
- Assume we limit our x -values: $-5 \leq x \leq 5$.
- In that domain we want to sample functions that are reasonably *smooth*.
- We use a *covariance matrix* to ensure that points *close together* in input space will produce output values that are also close together.




Gaussian processes

- A Gaussian process defines a prior over functions, which, given observed data, can be converted into a posterior.
- Instead of explicitly representing a distribution over a function, choose a finite set of points $\{x_1, \dots, x_n\}$.
- A GP assumes that the function values $f(x_1), \dots, f(x_n)$ has a **jointly Gaussian distribution** with some mean $\mu(\mathbf{x})$ and covariance $\Sigma(\mathbf{x})$, given by

$$\Sigma_{ij} = k(x_i, x_j)$$

where k is a positive definite **kernel**.




Multivariate normal distribution

- Gaussian processes are based on the **multivariate normal (Gaussian) distribution**

$$\mathbf{X} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{|2\pi\boldsymbol{\Sigma}|}} \exp \left[-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right]$$

where $\boldsymbol{\Sigma}$ is a **covariance matrix**.

- The covariance matrix determines the shape of the "bell"




Multivariate normal distribution

- Viewing from above, if the contours form a perfect circle, the variables are independent, and the covariance is zero

$$\Sigma = \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

- With covariance $\Sigma_{12} \neq 0$, the contour will have a more oval shape.



Gaussian processes

- Assume that we want to learn a function f from data $\mathcal{D} = \{(x_i, y_i): i = 1, \dots, D\}$
- Assume we have a distribution $p(f)$ over functions,
- Now: $p(f)$ is a Gaussian process if for a finite subset $\{x_1, \dots, x_n\}$ the marginal distribution over that subset

$$p(f(x_1), \dots, f(x_n))$$

has a Gaussian distribution.

- Now, if the prior is Gaussian, so is the posterior

$$p(f|\mathcal{D}) = \frac{p(\mathcal{D}|f)p(f)}{p(\mathcal{D})}$$

Gaussian processes

- The Gaussian process is parameterized by a mean vector and a covariance matrix

$$p \begin{pmatrix} f(x_1) \\ f(x_2) \end{pmatrix} \sim N(\boldsymbol{\mu}(\mathbf{x}), \boldsymbol{\Sigma})$$

where

$$\boldsymbol{\mu}(\mathbf{x}) = \begin{bmatrix} \mu(x_1) \\ \mu(x_2) \end{bmatrix}, \quad \boldsymbol{\Sigma} = \begin{bmatrix} k(x_1, x_1) & k(x_1, x_2) \\ k(x_2, x_1) & k(x_2, x_2) \end{bmatrix}$$

and $k(x_1, x_2)$ is a kernel function.

Gaussian processes

- So: we have the following data

$$\mathcal{D} = \{(x_i, y_i) : i = 1, \dots, D\}$$

and for some new point x^* we want to predict y^* .

- To do this we want to find a function f such that


$$y_i = f(x_i)$$

- Instead we assume that $\{f(x_1), \dots, f(x_D)\}$ follow a joint normal distribution, and use it to compute the posterior distribution

$$p(f(x^*) | x^*, \mathcal{D}, f) = N(\mu^*, \Sigma^*)$$


Gaussian processes

- Intuitively, we begin by sampling from our prior distribution




Gaussian processes

- Intuitively, we begin by sampling from our prior distribution
- We use our training data to represent the outputs of the unknown function.




Gaussian processes

- Intuitively, we begin by sampling from our prior distribution
- We use our training data to represent the outputs of the unknown function.
- And update the posterior.

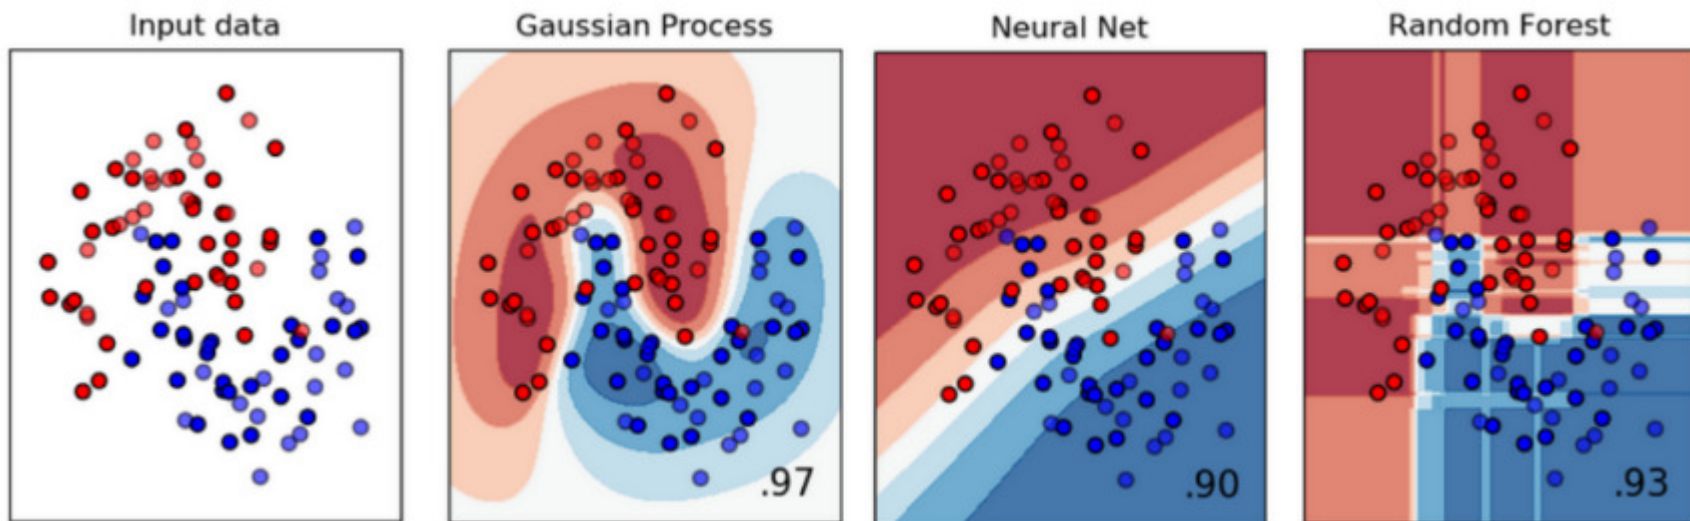


Gaussian processes

- Where does the kernel k come from?
- The covariance matrix characterizes the similarities between nearby points.
- The same range of kernels available as for SVMs.
 - The squared exponential (SE)
 - The radial basis function (RBF)
 - The Matern
 - ...



Gaussian processes



Advantages of GPs over SVMs

- GPs handle uncertainty in unknown function f by averaging, not minimizing
- GPs can learn kernel parameters from data, no matter how flexible we want to make the kernel
- GPs can learn regularization parameters without cross-validation.
- Can incorporate interpretable noise models and priors over functions, and can sample from prior to get intuitions about the model assumptions.
- We can combine automatic feature selection with learning using Automatic Relevance Determination (ARD)

